



Deal Qualification Executive Summary

Primary Opportunity

Tell us about your compute needs

Share what you're looking for — GPU model and quantity, workload type (training, inference, HPC, etc.), timeline, budget range, number of distinct deployments or end users, and any special requirements for underwriting, contracting, or delivery.

Cluster Specifications

GPU Model	<i>e.g., NVIDIA B300 (HGX Platform)</i>
Nodes per Cluster	<i>e.g., 128 nodes per cluster</i>
GPUs per Cluster	<i>e.g., 1,024 GPUs per cluster (8x GPUs per node)</i>
Cluster Configuration	<i>e.g., node count, CPUs, RAM, storage, networking/DPU hardware per node</i>
Initial Power per Cluster	<i>e.g., 2.5 MW IT load per cluster</i>
Scale Potential per Cluster	<i>e.g., scalable — additional capacity available</i>
Number of Clusters	<i>e.g., 2 — distinct, independent deployments</i>
Total Combined GPUs at Launch	<i>e.g., 2,048 GPUs (2 × 1,024)</i>
Total Combined Scale Potential	<i>e.g., TBD — expansion path to be confirmed</i>
Co-location Requirement	<i>e.g., enterprise data center; full cage / dedicated suite required</i>
Networking	<i>e.g., InfiniBand or 200–400GbE Ethernet</i>
Cooling	<i>e.g., air-cooled or DLC acceptable</i>
Cluster Co-location	<i>e.g., may be co-located or at separate facilities</i>

Commercial Terms

Lease Structure	<i>e.g., 100% reserved capacity, take-or-pay</i>
Target Price Range	<i>e.g., \$3.70–\$3.75 / GPU / hour</i>
Down Payment	<i>e.g., 20–25% of total contract value</i>
Estimated Contract Value	<i>e.g., total contract value over the term</i>

Deployment Model	<i>e.g., lease, purchase, or hybrid; preferred term length</i>
Target RFS Date	<i>e.g., within 90–120 days of contract execution</i>

Buyer Profile and Qualification

End User Description	<i>e.g., industry, use case, and level of confidentiality required</i>
Buyer Type	<i>e.g., enterprise / institutional; direct end user or reseller</i>
U.S. Entity	<i>e.g., yes/no; deployment jurisdiction</i>
Proof of Funds / Credit	<i>e.g., willingness and timeline to provide POF</i>
Budget Status	<i>e.g., approved and allocated / pending approval</i>
Take-or-Pay Willingness	<i>e.g., yes/no; comfort with reserved take-or-pay structure</i>
Workload Type	<i>e.g., AI/ML training, inference, HPC</i>
Deployment Type	<i>e.g., lease / GPUaaS / hardware purchase; preferred term</i>

Deployment Geography

Primary Preference	<i>e.g., United States — domestic deployment required</i>
Acceptable Locations	<i>e.g., preferred regions/markets and any that are excluded</i>
Cluster Separation Requirements	<i>e.g., any requirement to keep clusters co-located or separate</i>